Ashwin Sarathi Krishnan   Student ID: 200542907

NC State University

Department of Electrical and Computer Engineering

ECE 463/563 (Prof. Rotenberg)

Project #1: Cache Design, Memory Hierarchy Design

REPORT TEMPLATE (Version 1.0)

by

Ashwin Sarathi Krishnan

NCSU Honor Pledge: "I have neither given nor received unauthorized aid on this project."

Student's electronic signature: Ashwin Sarathi Krishnan

Course number: 563

Report template for Project #1.

1. L1 cache exploration: SIZE and ASSOC

GRAPH #1   (*total number of simulations: 55*)

For this experiment:
- Benchmark trace: gcc_trace.txt
- L1 cache: SIZE is varied, ASSOC is varied, BLOCKSIZE = 32.
- L2 cache: None.
- Prefetching: None.

Plot L1 miss rate on the y-axis versus $\log_2(\text{SIZE})$ on the x-axis, for eleven different cache sizes: SIZE = 1KB, 2KB, … , 1MB, in powers-of-two. (That is, $\log_2(\text{SIZE})$ = 10, 11, …, 20.) The graph should contain five separate curves (*i.e.*, lines connecting points), one for each of the following associativities: direct-mapped, 2-way set-associative, 4-way set-associative, 8-way set-associative, and fully-associative. All points for direct-mapped caches should be connected with a line, all points for 2-way set-associative caches should be connected with a line, *etc*.
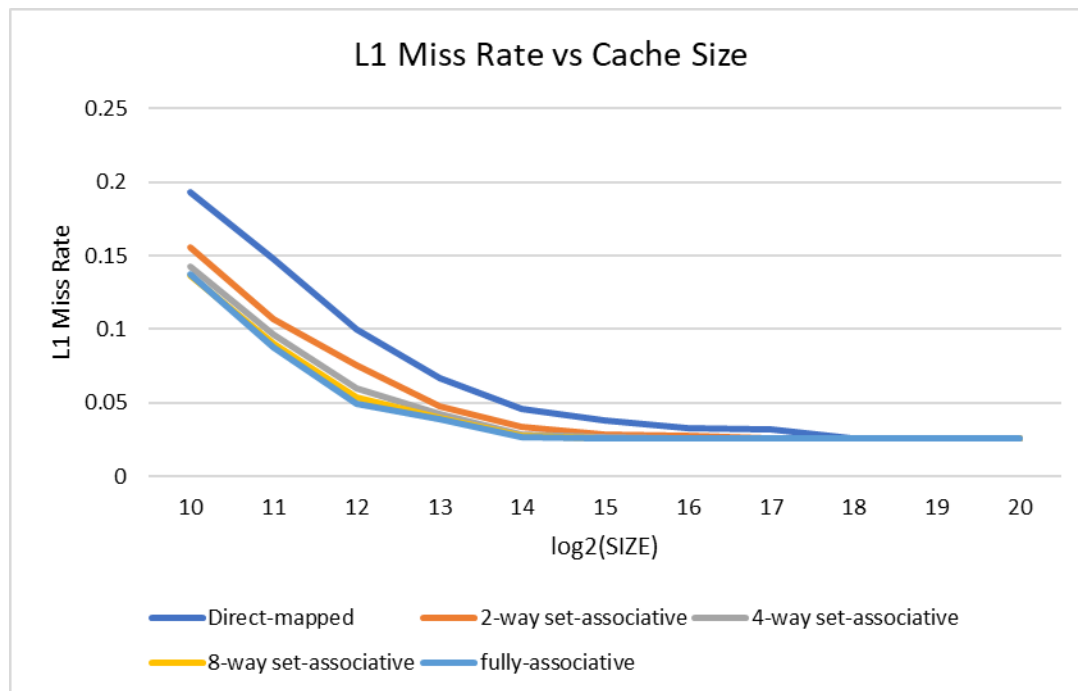
Answer the following questions:
1. For a given associativity, how does increasing cache size affect miss rate?

From the graph above, we can conclude that for a given associativity and block size, as cache size increases, miss rate decreases, until there is only compulsory miss rate left.

Report template for Project #1.

2. For a given cache size, how does increasing associativity affect miss rate?

> **From the graph, we can conclude that for a given cache size and block size, as associativity increases, miss rate decreases, until there is only compulsory miss rate left.**

3. Estimate the *compulsory miss rate* from the graph and briefly explain how you arrived at this estimate.

> compulsory miss rate = **0.0258**
> How I arrived at this estimate: **The absolute minimum miss rate for a given block size when varying the cache size and associativity is the compulsory miss rate.**

Report template for Project #1.

GRAPH #2   (*no additional simulations with respect to GRAPH #1*)

Same as GRAPH #1, but the y-axis should be AAT instead of L1 miss rate.
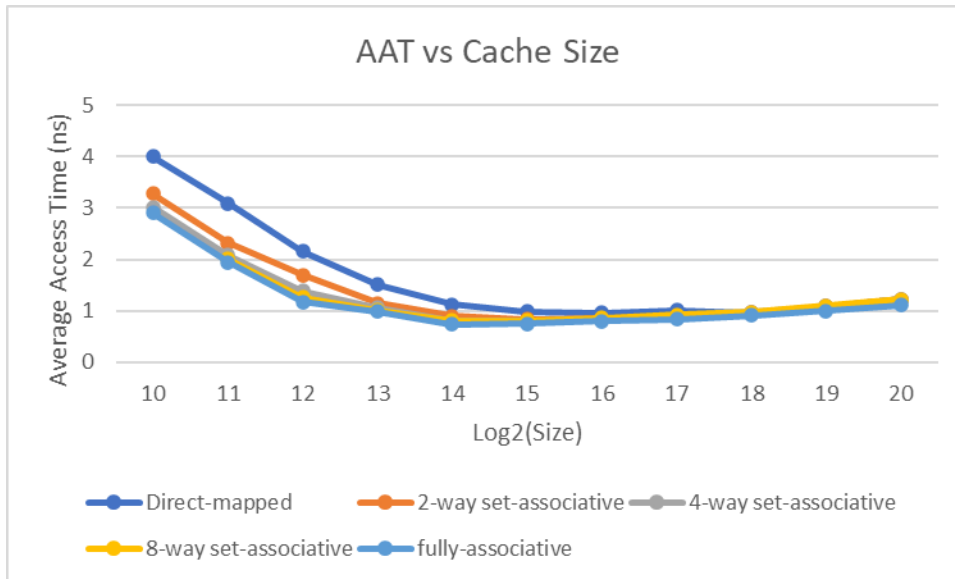
Answer the following question:
1. For a memory hierarchy with only an L1 cache and BLOCKSIZE = 32, which configuration yields the best (*i.e.*, lowest) AAT and what is that AAT?
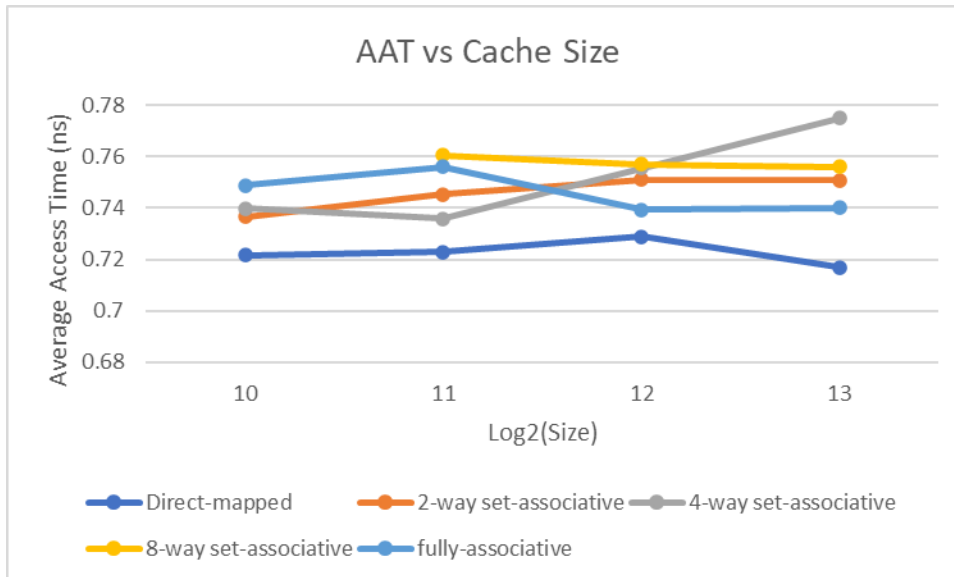
Configuration that yields the lowest AAT: **16KB cache, fully associative**
Lowest AAT: **0.7342**

Report template for Project #1.

GRAPH #3   (*total number of simulations: 20*)

Same as GRAPH #2, except make the following changes:
- Add the following L2 cache to the memory hierarchy: 16KB, 8-way set-associative, same block size as L1 cache.
- Vary the L1 cache size only between 1KB and 8KB (since L2 cache is 16KB).

Answer the following questions:
1. With the L2 cache added to the system, which L1 cache configuration yields the best (*i.e.*, lowest) AAT and what is that AAT?

L1 configuration that yields the lowest AAT with 16KB 8-way L2 added: **8KB, direct mapped**
Lowest AAT: **0.717**

2. How does the lowest AAT with L2 cache (GRAPH #3) compare with the lowest AAT without L2 cache (GRAPH #2)?

The lowest AAT with L2 cache is **0.0172 ns less than** the lowest AAT without L2 cache.

3. Compare the *total area* required for the lowest-AAT configurations with L2 cache (GRAPH #3) versus without L2 cache (GRAPH #2).

Total area for lowest-AAT configuration with L2 cache =
     0.053293238 mm^2 (L1 area) + **0.130444675** mm^2 (L2 area) = **0.183737913** mm^2 (total area)

Total area for lowest-AAT configuration without L2 cache = **0.063446019** mm^2 (L1 area)

The total area of the lowest-AAT configuration with L2 cache is **189.59723 % more than** the total area of the lowest-AAT configuration without L2 cache.
FYI: How to calculate % difference of x with respect to y:
If x > y: x is ((x-y)/y * 100%) more than y.

Report template for Project #1.

If x < y: x is ((y-x)/y * 100%) less than y.

2. L1 cache exploration: SIZE and BLOCKSIZE

GRAPH #4   (*total number of simulations: 24*)

For this experiment:
- Benchmark trace: gcc_trace.txt
- L1 cache: SIZE is varied, BLOCKSIZE is varied, ASSOC = 4.
- L2 cache: None.
- Prefetching: None

Plot L1 miss rate on the y-axis versus $\log_2$(BLOCKSIZE) on the x-axis, for four different block sizes: BLOCKSIZE = 16, 32, 64, and 128. (That is, $\log_2$(BLOCKSIZE) = 4, 5, 6, and 7.) The graph should contain six separate curves (*i.e.*, lines connecting points), one for each of the following L1 cache sizes: SIZE = 1KB, 2KB, …, 32KB, in powers-of-two. All points for SIZE = 1KB should be connected with a line, all points for SIZE = 2KB should be connected with a line, *etc*.
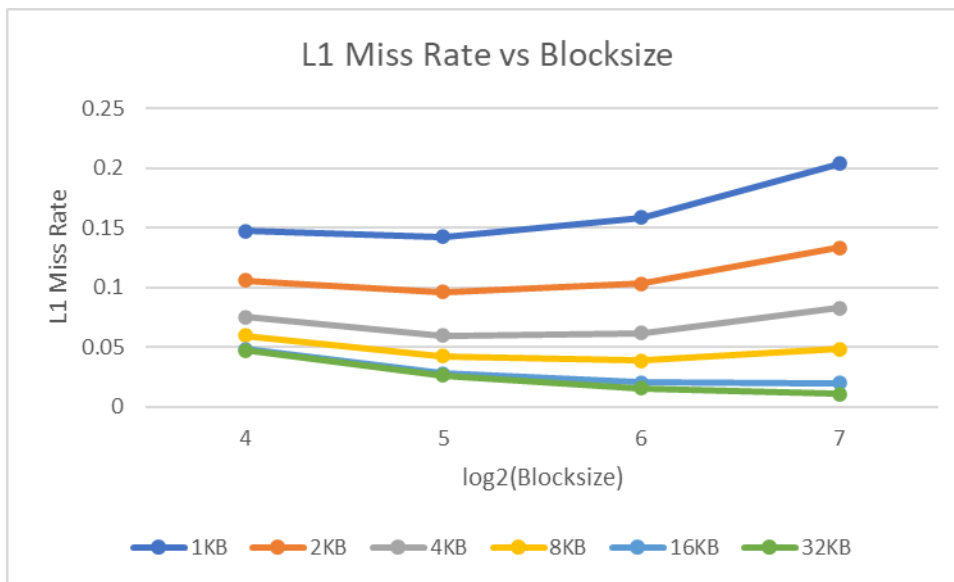
Answer the following questions:
1. Do smaller caches prefer smaller or larger block sizes?

Smaller caches prefer **smaller** block sizes. For example, the smallest cache considered in Graph #4 (1KB) achieves its lowest miss rate with a block size of **16 B**.

2. Do larger caches prefer smaller or larger block sizes?

Larger caches prefer **larger** block sizes. For example, the largest cache considered in Graph #4 (32KB) achieves its lowest miss rate with a block size of **32 B**.

3. As block size is increased from 16 to 128, is the tension between *exploiting more spatial locality* and *cache pollution* evident in the graph? Explain.

Yes, the tension between *exploiting more spatial locality* and *cache pollution* is evident in the graph.

Report template for Project #1.

For example, consider the smallest (1KB) cache in Graph #4. Increasing block size from **16 B to 32 B** is helpful (reduces miss rate) due to **exploiting more spatial locality**. But then increasing block size further, from **32 B to 64 B**, is not helpful (increases miss rate) due to **cache pollution** having greater effect.

3. L1 + L2 co-exploration

GRAPH #5   (*total number of simulations: 12*)

For this experiment:
- Benchmark trace: gcc_trace.txt
- L1 cache: SIZE is varied, BLOCKSIZE = 32, ASSOC = 4.
- L2 cache: SIZE is varied, BLOCKSIZE = 32, ASSOC = 8.
- Prefetching: None.

Plot AAT on the y-axis versus $\log_2$(L1 SIZE) on the x-axis, for four different L1 cache sizes: L1 SIZE = 1KB, 2KB, 4KB, 8KB. (That is, $\log_2$(L1 SIZE) = 10, 11, 12, 13.) The graph should contain three separate curves (*i.e.*, lines connecting points), one for each of the following L2 cache sizes: 16KB, 32KB, 64KB. All points for the 16KB L2 cache should be connected with a line, all points for the 32KB L2 cache should be connected with a line, *etc*.
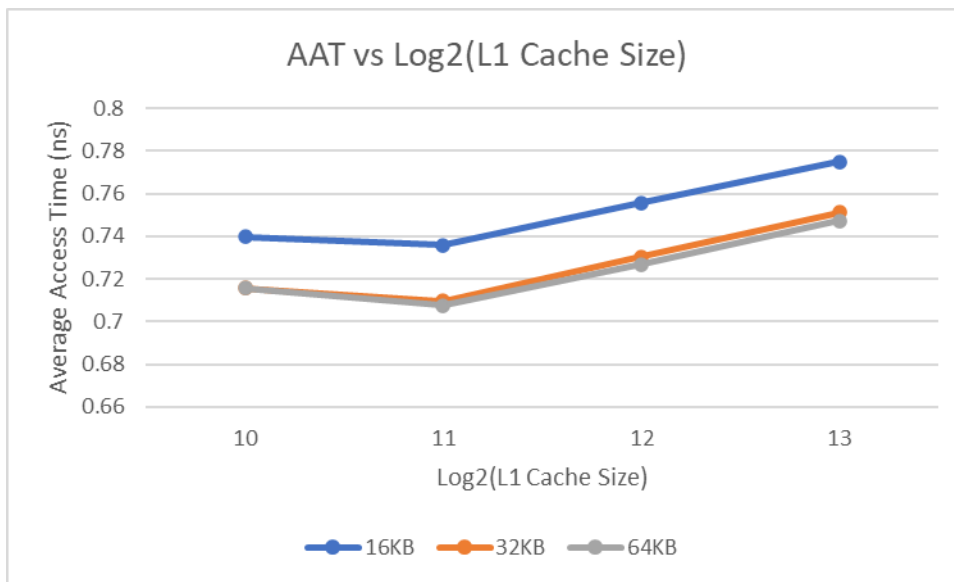
Answer the following question:
1. Which memory hierarchy configuration in Graph #5 yields the best (*i.e.*, lowest) AAT and what is that AAT?

Configuration that yields the lowest AAT: **2KB L1 and 64KB L2**, where L1 is 4-way set associative and L2 is 8-way set-associative
Lowest AAT: **0.7077**

Report template for Project #1.

4. Stream buffers study (ECE 563 students only)

TABLE #1   (*total number of simulations: 5*)

For this experiment:
- Microbenchmark: stream_trace.txt
- L1 cache: SIZE = 1KB, ASSOC = 1, BLOCKSIZE = 16.
- L2 cache: None.
- PREF_N (number of stream buffers): 0 (pref. disabled), 1, 2, 3, 4
- PREF_M (number of blocks in each stream buffer): 4

The trace "stream_trace.txt" was generated from the loads and stores in the loop of interest of the following microbenchmark:

```
#define SIZE 1000

uint32_t a[SIZE];
uint32_t b[SIZE];
uint32_t c[SIZE];

int main(int argc, char *argv[]) {
…
   // LOOP OF INTEREST
   for (int i = 0; i < SIZE; i++)
      c[i] = a[i] + b[i];    // per iteration: 2 loads (a[i], b[i]) and 1 store (c[i] = …)
…
}
```

Fill in the following table and answer the following questions:

PREF_N, PREF_M	L1 miss rate
0,0 (pref. disabled)	
1,4	
2,4	
3,4	
4,4	

1. For this streaming microbenchmark, with prefetching disabled, do L1 cache size and/or associativity affect the L1 miss rate (feel free to simulate L1 configurations besides the one used for the table)? Why or why not?

With prefetching disabled, L1 cache size and/or associativity << do / do not >> affect L1 miss rate (for this streaming microbenchmark).

The reason: _____.

2. For this streaming microbenchmark, what is the L1 miss rate with prefetching disabled? Why is it that value, *i.e.*, what is causing it to be that value? Hint: each element of arrays a, b, and c, is 4 bytes (uint32_t).

The L1 miss rate with prefetching disabled is _____, because

_____.

Report template for Project #1.

3. For this streaming microbenchmark, with prefetching disabled, what would the L1 miss rate be if you doubled the block size from 16B to 32B? (hypothesize what it will be and then check your hypothesis with a simulation)

The L1 miss rate with prefetching disabled and a block size of 32B is _____, because _____.

4. With prefetching enabled, what is the minimum number of stream buffers required to have any effect on L1 miss rate? What is the effect on L1 miss rate when this many stream buffers are used: specifically, is it a modest effect or huge effect? Why are this many stream buffers required? Why is using fewer stream buffers futile? Why is using more stream buffers wasteful?

Minimum number of stream buffers needed to have any effect on L1 miss rate: _____

With this many stream buffers, the effect on L1 miss rate is << modest / huge >>. Specifically, the L1 miss rate is nearly _____. We only miss on the first elements of _____ (hence a total of _____ misses).

This many stream buffers are required because _____.

Using fewer stream buffers is futile because _____.

Using more stream buffers is wasteful because _____.

Report template for Project #1.